

Import the libraries as shown

```
import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
```

```
# re-size all the images to this
IMAGE_SIZE = [185, 500]
BATCH_SIZE = 32
CHANNELS = 3
EPOCHS = 10
```

```
directory_path = ('/content/drive/MyDrive/data/data/train')
```

```
dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory_path,
    shuffle=True,
    image_size=(185, 500),
    batch_size=32
)
```

```
Found 4188 files belonging to 4 classes.
```

```
class_names = dataset.class_names
class_names
```

```
['Blight', 'Common_Rust', 'Gray_Leaf_Spot', 'Healthy']
```

```
len(dataset)
```

```
131
```

```
plt.figure(figsize=(10, 10))
for image_batch, labels_batch in dataset.take(1):
    for i in range(12):
        ax = plt.subplot(3, 4, i + 1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")
```



Blight



Healthy



Healthy



Common_Rust



Healthy



Gray_Leaf_Spot



Healthy



Gray_Leaf_Spot




Healthy




Common_Rust



Common_Rust



Common_Rust



```
len(dataset)
```



131

```
# 80% ==> training
# 20% ==>validation,10% test
```

```
train_size = 0.8
len(dataset)*train_size
```



104.80000000000001

```
train_ds = dataset.take(104)
len(train_ds)
```



104

```
test_ds = dataset.skip(104)
len(test_ds)
```



27

```
val_size = 0.1
len(dataset)*val_size
```



13.100000000000001

```
val_ds = test_ds.take(13)
len(val_ds)
```



13

```
test_ds = test_ds.skip(13)
len(test_ds)
```



14

```
def get_dataset_partition_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1, shuffle=True, shuffle_size=10000):
    ds_size = len(ds)

    if shuffle:
        ds = ds.shuffle(shuffle_size, seed=12)
    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)

    train_ds = ds.take(train_size)

    val_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds, val_ds, test_ds
```

```
train_ds, val_ds, test_ds = get_dataset_partition_tf(dataset)
```

```
len(train_ds)
```

```
↩ 104
```

```
len(val_ds)
```

```
↩ 13
```

```
len(test_ds)
```

```
↩ 14
```

```
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
```


```
from tensorflow.keras.layers import Resizing, Rescaling
```

```
resize_and_rescale = tf.keras.Sequential([
    Resizing(IMAGE_SIZE, IMAGE_SIZE),
    Rescaling(1./255),
])
```

```
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
])
```

```
# Build the model
model = models.Sequential([
    layers.Input(shape=(IMAGE_SIZE[0], IMAGE_SIZE[1], CHANNELS)),
    data_augmentation,
    layers.Rescaling(1./255),
    layers.Conv2D(32, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dense(len(class_names), activation='softmax')
])
```

```
# Model summary
model.summary()
```

 Model: "sequential_2"

Layer (type)	Output Shape	Param #
sequential_1 (Sequential)	(None , 185 , 500 , 3)	0
rescaling_1 (Rescaling)	(None , 185 , 500 , 3)	0
conv2d (Conv2D)	(None , 183 , 498 , 32)	896
max_pooling2d (MaxPooling2D)	(None , 91 , 249 , 32)	0
conv2d_1 (Conv2D)	(None , 89 , 247 , 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None , 44 , 123 , 64)	0
conv2d_2 (Conv2D)	(None , 42 , 121 , 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None , 21 , 60 , 128)	0
conv2d_3 (Conv2D)	(None , 19 , 58 , 128)	147,584
max_pooling2d_3 (MaxPooling2D)	(None , 9 , 29 , 128)	0
flatten (Flatten)	(None , 33408)	0
dense (Dense)	(None , 512)	17,105,408
dense_1 (Dense)	(None , 4)	2,052

Total params: [17,348,292](#) (66.18 MB)
Trainable params: [17,348,292](#) (66.18 MB)
Non-trainable params: [0](#) (0.00 B)


```
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)

from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam

model.compile(optimizer=Adam(), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

# Train the model
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=10,
    callbacks=[early_stopping]
)
```


 Epoch 1/10

104/104	925s	9s/step	- accuracy: 0.7890	- loss: 0.5288	- val_accuracy: 0.8606	- val_loss: 0.3
Epoch 2/10						
104/104	965s	9s/step	- accuracy: 0.8191	- loss: 0.4317	- val_accuracy: 0.8558	- val_loss: 0.4
Epoch 3/10						
104/104	898s	8s/step	- accuracy: 0.8214	- loss: 0.4421	- val_accuracy: 0.8582	- val_loss: 0.3
Epoch 4/10						
104/104	934s	9s/step	- accuracy: 0.8149	- loss: 0.4282	- val_accuracy: 0.8726	- val_loss: 0.3
Epoch 5/10						
104/104	881s	8s/step	- accuracy: 0.8531	- loss: 0.3505	- val_accuracy: 0.7260	- val_loss: 0.8
Epoch 6/10						
104/104	890s	9s/step	- accuracy: 0.8529	- loss: 0.3597	- val_accuracy: 0.8750	- val_loss: 0.3
Epoch 7/10						
104/104	878s	8s/step	- accuracy: 0.8583	- loss: 0.3321	- val_accuracy: 0.8870	- val_loss: 0.2
Epoch 8/10						
104/104	877s	8s/step	- accuracy: 0.8822	- loss: 0.3119	- val_accuracy: 0.8750	- val_loss: 0.2
Epoch 9/10						
104/104	926s	8s/step	- accuracy: 0.8573	- loss: 0.3596	- val_accuracy: 0.8942	- val_loss: 0.2
Epoch 10/10						
104/104	916s	8s/step	- accuracy: 0.8663	- loss: 0.3386	- val_accuracy: 0.8894	- val_loss: 0.3

```
scores = model.evaluate(test_ds)
```

 **14/14**  **60s** 2s/step - accuracy: 0.8678 - loss: 0.2834

```
scores
```

 [0.2650536596775055, 0.8816964030265808]

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix

# Function to get predictions and true labels
def get_predictions_and_labels(ds):
    images, labels = [], []
    for image_batch, label_batch in ds:
        images.extend(image_batch.numpy())
        labels.extend(label_batch.numpy())
    return np.array(images), np.array(labels)

# Get test images and labels
test_images, test_labels = get_predictions_and_labels(test_ds)

# Make predictions
predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

# Generate confusion matrix
cm = confusion_matrix(test_labels, predicted_labels)


# Plot confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

# Generate and print classification report
report = classification_report(test_labels, predicted_labels, target_names=class_names)
print(report)
```

 14/14

33s

2s/step



```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(range(10), acc, label='Training Accuracy')
plt.plot(range(10), val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(range(10), loss, label='Training Loss')
plt.plot(range(10), val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```




```
import numpy as np
for images_batch, labels_batch in test_ds.take(1):
    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:",class_names[first_label])

    batch_prediction = model.predict(images_batch)
    print("predicted label:",class_names[np.argmax(batch_prediction[0])])
```



first image to predict
actual label: Gray_Leaf_Spot
1/1 ————— 3s 3s/step
predicted label: Gray_Leaf_Spot













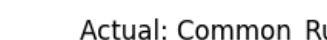
```
def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array, 0)
    predictions = model.predict(img_array)
    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence
```

```
plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))

        predicted_class, confidence = predict(model, images[i].numpy())
        actual_class = class_names[labels[i]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confidence: {confidence}%")
        plt.axis("off")
```

1/1		0s 125ms/step
1/1		0s 128ms/step
1/1		0s 140ms/step
1/1		0s 132ms/step
1/1		0s 132ms/step
1/1		0s 133ms/step
1/1		0s 131ms/step
1/1		0s 131ms/step
1/1		0s 134ms/step

